

The LC-3

Abstract

- The ISA of LC-3
 - Memory Organization
 - Registers
 - Instruction Set
 - Operate
 - Data Movement
 - Control
- Microarchitecture of LC-3. (Datapath)

Tip

It's worth noting that in the notes I skip many examples on the book. However, if you have not taken the course *Assembly Language*, or not familiar with programming in assembly language, please read the example on the textbook carefully and practice it sufficiently.

1 The ISA: Overview

The ISA specifies the *memory organization*, *register set* and *instruction set*, including the opcodes, data types, addressing modes of the instruction in the instruction set.

1.1 Memory Organization

- **address space:** 2^{16} (i.e. 65536) locations.
Not all 65536 addresses are actually used for memory locations.
- **addressability:** 16 bits. We refer to 16 bits as 1 **word** in the LC-3, so it's also called **word-addressable**.

1.2 Registers

As mentioned before, each register is called a **GPR (General Purpose Register)**. They are referred to as $R0, R1, \dots R7$.

Registers store information that can be operated on later, and they are stored register files, where CPU can access them faster than memory.

1.3 The Instruction Set

An instruction is made up of its opcode and operands.

The instruction set is defined by its set of opcodes, data types and addressing modes.

1.3.1 Opcodes

The LC-3 ISA has 15 instructions and the opcode is specified in `bit[15:12]`. The code 1101 has been left unspecified.

1.3.2 Data Types

Every opcode will interpret the bit patterns of its operands according to the data type it's designed to support.

The same bit pattern can correspond to different number, depending on instructions that interpret it. *e.g.* `ADD R2, R1, #1`, `LD R2, #1`. In `ADD`, `#1` is interpreted as a number for arithmetic, while in `LD`, `#1` is interpreted as an address.

1.3.3 Addressing Modes

An operand can generally be found in one of 3 places: *in memory*, *in a register* or *as a part of the instruction*. If the operand of the instruction, we refer to it as a **literal** or as an **immediate** operand.

The LC-3 supports 5 addressing modes: *immediate* (or *literal*), *register* and 3 memory addressing modes: *PC-relative*, *indirect* and *Base+offset*.

2 The Instruction Set of the LC-3

2.1 Operate Instructions

2.1.1 NOT

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	0	0	1	0	1	1	1	0	1	1	1	1	1	1	1
NOT				R3			R5								

The **NOT** (opcode=1001) instruction is the only operate instruction that performs a *unary* operation. `bits[11:9]` is DR, `bits[8:6]` is SR, and `bits[5:0]` are set to 1.

The datapath of NOT is as follows:

chapter 5 The LC-3

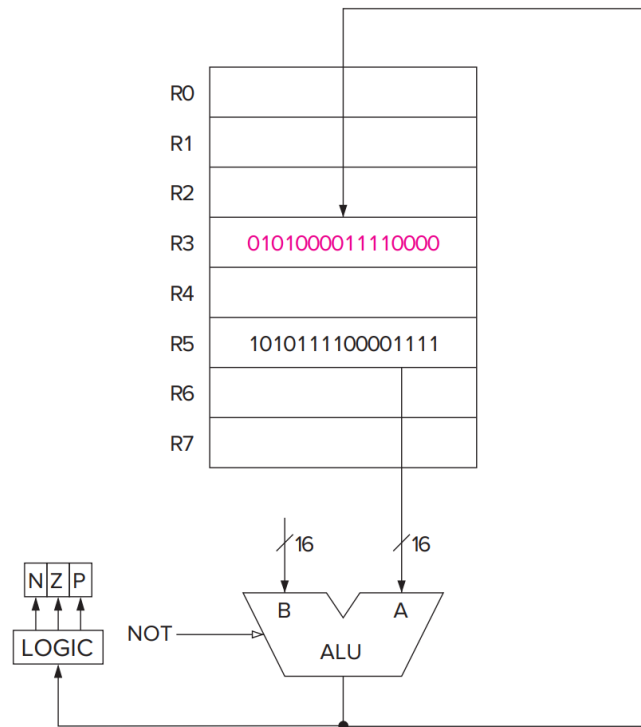


Figure 5.4 Data path relevant to the execution of NOT R3, R5.

For binary operations (like ADD, AND) the datapath is almost the same except the operand B which is also from the register files.

2.1.2 Immediates In Operate Instructions

The diagram below shows `ADD R1, R4, #-2`. Note that we the operand B is computed by *signed-extending* `bits[4:0]` to 16 bits. And since there is only 5 bits to store immediates, not all 2's complement intergers can be immediate operands. (only $[-16, 15]$)

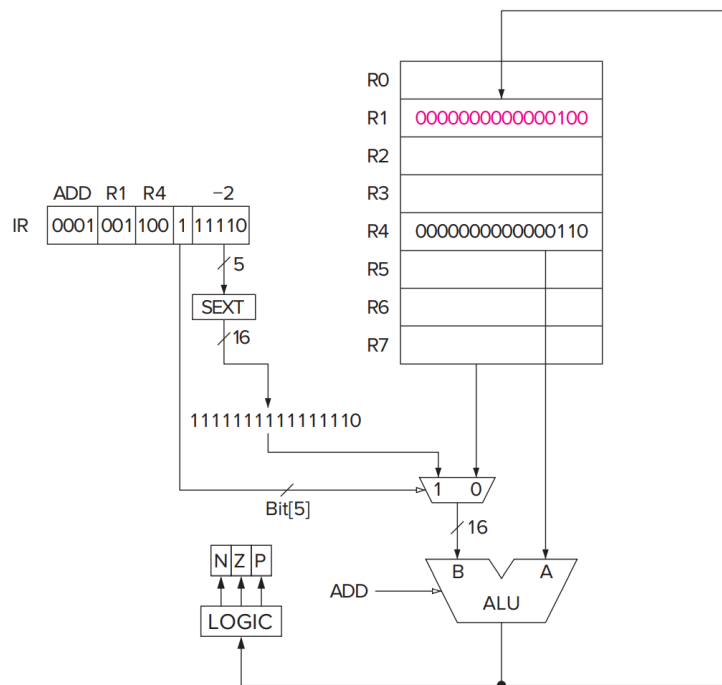


Figure 5.5 Data path relevant to the execution of `ADD R1, R4, #-2`.

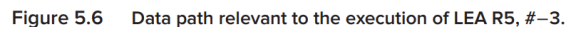
With the help of NOT and ADD (with immediate), we can implement the *subtraction*. (Recall that in Chapter 2, the negative of an integer represented in 2's complement can be obtained by complementing the number and adding 1)

2.1.3 The LEA Instruction

LEA (opcode=1110) loads the register specified by `bits[11:9]` of the instruction with the value formed by adding the incremented PC to the *sign-extended* `bits[8:0]` of the instruction.


LEA is useful to *initialize a register with an address*.

Note that the values to be loaded into the register *does not involve any access to memory*, and it does not affect CC.



The process of moving information from memory to a register is called **load**, and the process of moving information from a register to a memory is called **store**.

The format:



- **Base+offset Mode: LDR and STR**

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	1	1	0	0	0	1	0	1	0	0	1	1	1	0	1
LDR				R1			R2			x1D					

Note that load instructions will influence CC, while store not. The value that is finally read from memory and will be loaded to the register determine CC.

It's also worth noting that if some bits specify DR in load instructions, then it will specify SR in store instructions.

2.2.1 PC-Relative Mode

LD (opcode=0010) and **ST** (opcode=0011) specify the **PC-relative** addressing mode.

The memory address is computed by *signed-extending* **bits[8:0]** to 16 bits and adding the result to the *incremented PC* (incremented during FETCH phase).

Note that the address of the memory operand is limited to a small range. ($[-255, +256]$)

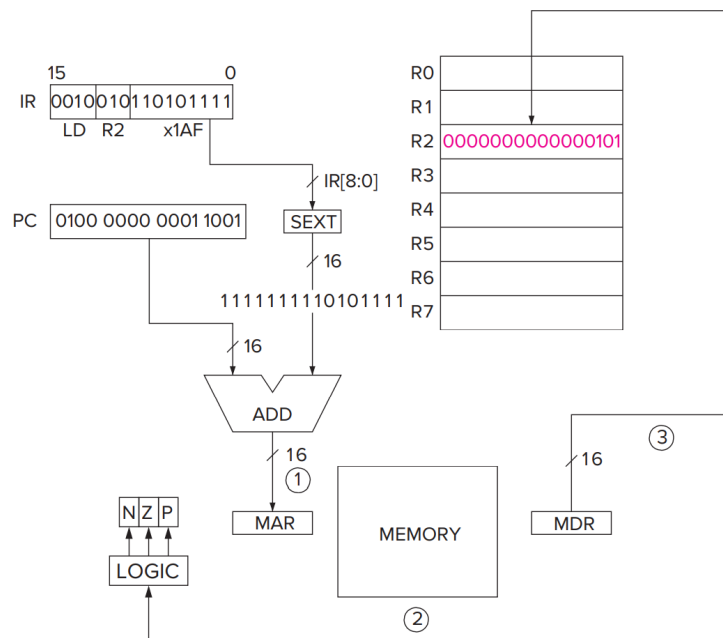


Figure 5.7 Data path relevant to execution of LD R2, x1AF.

2.2.2 Indirect Mode

LDI (opcode=1010) and **STI**(opcode=1011) specify the **indirect** addressing mode.

An address is first formed exactly the same way as with LD, however, the result is the address of the address of the operand. So we need to interrogate memory twice.

Note that the address of the operand can be *anywhere*.

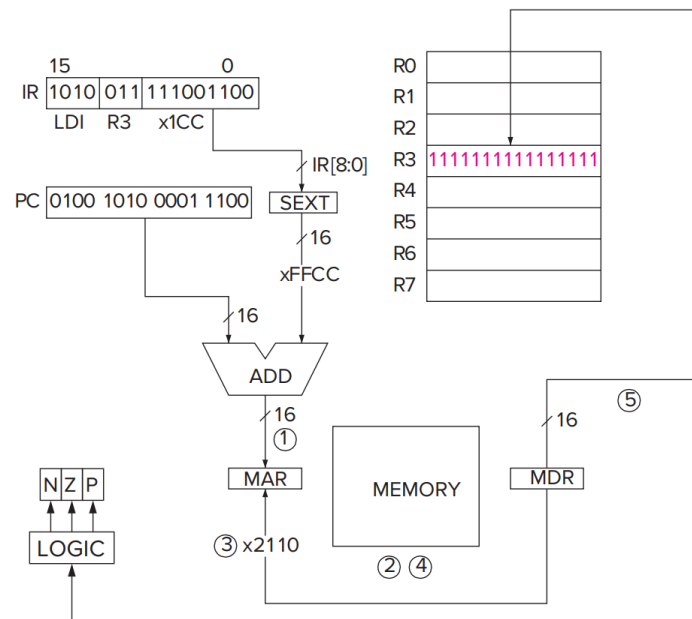


Figure 5.8 Data path relevant to the execution of LDI R3, x1CC.

In the example above, the incremented PC is **x4A1C** and the sign-extended offset is **xFFCC**. So we first get the address **x49E8** and get the data in **x49E8**, which is **x2110**. Then we obtain the data in **x2110** and load the value into the register **R3**.

2.2.3 Base+offset Mode

LDR (opcode=0110) and **STR** (opcode=0111) specify the **Base+offset** addressing mode.

The address is obtained by adding a *signed-extended* 6-bit(**bits[5:0]**) offset to a base register(**bits[8:6]**).

Note that the address of the operand can also be *anywhere*.

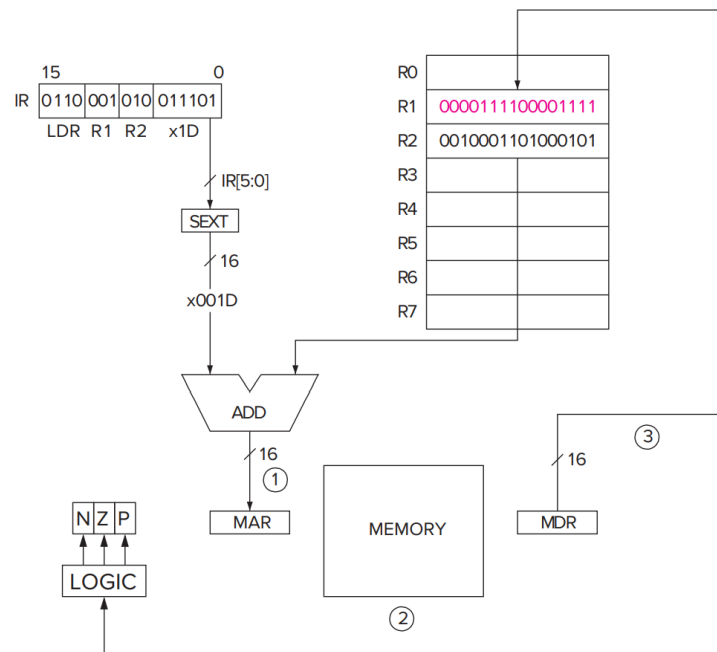


Figure 5.9 Data path relevant to the execution of LDR R1, R2, x1D.

2.3 Control Instructions

The LC-3 has 5 opcodes that enable the sequential execution flow to be broken: *conditional branch*, *unconditional jump*, *subroutine call*(funtion), *TRAP* (service call) and *RTI*.

2.3.1 Conditional Branches

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	n	z	p	PCoffset								

BR (opcode=0000) uses condition codes to determine whether or not to depart from the usual sequential execution. Besides, not all condition codes will be inspected and we can determine which condition codes will be inspected. *e.g.* In `BRz x0D9`, `bit[11:9]=010` so we only check the condition code Z, which means if the result of the last instruction that can set CC is zero, then we will jump to the target address.

Additionally, if we request to check all condition codes, like `BRnzp x0D9`, it means an **unconditional branch**. Since the result of the operation can be positive, negative or zero, so at least one of CC will be set. Meanwhile, if `bits[11:9]=000`, nothing will happen, just like a nop (an instruction but do nothing and will cause no difference)

In summary, it is only *ADD*, *AND*, *LD*, *LDI*, *LDR*, *NOT* that will change CC after finishing operations.

2.3.2 The JMP Instruction

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	0	0	0	0	0	0	1	0	0	0	0	0	0	0
JMP								BaseR							

The **JMP** instruction (opcode=1100) loads the PC with the contents of the *register* specified by bits[8:6] of the instruction. (its addressing mode is by register)

2.3.3 The TRAP Instruction

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	0	0	0	0	trapvector							

The **TRAP** instruction (opcode=1111) changes the PC to a memory address that is part of the operating system so that the OS will perform some task on behalf of the program. Once the service call ends, the PC is set to be the address of the instruction following the TRAP instruction.

- Input a character from the keyboard (trapvector = x23)
- Output a character to the monitor (trapvector = x21)
- Halt the program (trapvector = x25)

Some useful trap service routines are given in the P675 Table A.3. The details of the TRAP instructions will be discussed in Chapter 9.

3 The Datapath

Here is only a figure of LC-3 datapath. Please be clear of the data flow for each instruction and the related control signals. A good way is to draw a datapath on your own hands.

